

## Local Search

## Search so far..

- A\*, BFS, DFS etc
  - Given set of states, get to goal state

7	2	4
5		6
8	3	1

Start State


	1	2
3	4	5
6	7	8

Goal State

- Need to know the *path* as well

## Example: $n$ -queens

- Put  $n$  queens on an  $n \times n$  board with no two queens on the same row, column, or diagonal



- How would you *represent* the state space of this problem?
- How is the problem different from the 8-puzzle?

## Local search algorithms

- In many optimization problems, the **path** to the goal is irrelevant; the goal state itself is the solution
- State space = set of "complete" configurations
- Find configuration satisfying constraints, e.g.,  $n$ -queens
- In such cases, we can use **local search algorithms**
- Keep a single "current" state, try to improve it
  - Assume access to a function,  $\text{Eval}(x)$  that tells you how good  $x$  is

## Hill-climbing search


- "Like climbing Everest in thick fog with amnesia"

```

function HILL-CLIMBING(problem) returns a state that is a local maximum
inputs: problem, a problem
local variables: current, a node
                 neighbor, a node
current ← MAKE-NODE(INITIAL-STATE[problem])
loop do
  neighbor ← a highest-valued successor of current
  if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
  current ← neighbor
    
```

## Hill-climbing search

- Problem: depending on initial state, can get stuck in local maxima




## Hill-climbing search: 8-queens problem

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	13	16	13	16	
17	14	17	15	14	16	16	
17	16	18	15	14	15	15	
18	14	15	15	14	14	16	
14	14	13	17	12	14	12	18

- $h$  = number of pairs of queens that are attacking each other, either directly or indirectly
- $h = 17$  for the above state

## Hill-climbing search: 8-queens problem



- A local minimum with  $h = 1$

## Simulated annealing search

- Idea: escape local maxima by allowing some "bad" moves but **gradually decrease** their frequency

```

function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
         schedule, a mapping from time to "temperature"
  local variables: current, a node
                 next, a node
                 T, a "temperature" controlling prob. of downward steps
  current ← MAKE-NODE(INITIAL-STATE[problem])
  for t ← 1 to ∞ do
    T ← schedule[t]
    if T = 0 then return current
    next ← a randomly selected successor of current
     $\Delta E \leftarrow \text{VALUE}[\text{next}] - \text{VALUE}[\text{current}]$ 
    if  $\Delta E > 0$  then current ← next
    else current ← next only with probability  $e^{\Delta E / T}$ 
  
```

## Properties of simulated annealing search


- One can prove: If  $T$  decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching 1
- Widely used in VLSI layout, airline scheduling, etc

## Local beam search

- Keep track of  $k$  states rather than just one
- Start with  $k$  randomly generated states
- At each iteration, all the successors of all  $k$  states are generated
- If any one is a goal state, stop; else select the  $k$  best successors from the complete list and repeat.

## Genetic algorithms

- Variant of local beam search with *sexual recombination*.



## Genetic Algorithms

- View optimization by analogy with evolutionary theory → Simulation of natural selection
- View configurations as *individuals* in a *population*
- View *Eval* as a measure of *fitness*
- Let the least-fit individuals die off without reproducing
- Allow individuals to *reproduce* with the best-fit ones selected more often
- Each *generation* should be overall better fit (higher value of *Eval*) than the previous one
- If we wait long enough the population should evolve so toward individuals with high fitness (i.e., maximum of *Eval*)

## Genetic Algorithms: Implementation

- Configurations represented by strings:

$$X = \begin{bmatrix} 1 & 0 & 0 & 1 & 1 & 0 & 0 & 1 \end{bmatrix}$$

- Analogy:
  - The string is the chromosome representing the individual
  - String made up of genes
  - Configuration of genes are passed on to offsprings
  - Configurations of genes that contribute to high fitness tend to survive in the population
- Start with a random population of  $P$  configurations and apply two operations
  - *Reproduction*: Choose 2 "parents" and produce 2 "offsprings"
  - *Mutation*: Choose a random entry in one (randomly selected) configuration and change it

## Genetic Algorithms: Reproduction

Parents:

1	0	0	1	1	0	0	1
1	0	1	1	0	0	0	1

## Genetic Algorithms: Reproduction

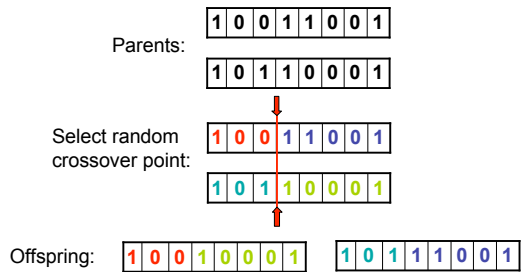
Parents:

1	0	0	1	1	0	0	1
1	0	1	1	0	0	0	1

Select random crossover point:

1	0	0	1	1	0	0	1
1	0	1	1	0	0	0	1


## Genetic Algorithms: Reproduction



- Each offspring receives part of the genes from each of the parents
- Implemented by crossover operation

## Genetic Algorithms: Mutation

- Random change of one element in one configuration
  - Implements random deviations from inherited traits
  - Corresponds loosely to “random walk”: Introduce random moves to avoid small local extrema



## Basic GA Outline

- Create initial population  $X = \{X_1, \dots, X_P\}$
  - Iterate:
    - 1. Select  $K$  random pairs of parents
    - 2. For each pair of parents  $p, q$ :
      - 1.1 Generate offsprings ( $Y_1, \dots, Y_r$ )
      - Variation: Generate only one offspring
      - Replace randomly selected element of the population by  $Y_i$
      - With probability  $\mu$ : Apply a random mutation to  $Y_i$
  - Return the best individual in the population
- Possible strategy: Select the best  $rP$  individuals ( $r < 1$ ) for reproduction and discard the rest → Implements selection of the fittest

## Genetic Algorithms: Selection

- Discard the least-fit individuals through threshold on  $Eval$  or fixed percentage of population
- Select best-fit (larger  $Eval$ ) parents in priority
- Example: Random selection of individual based on the probability distribution

$$\Pr(\text{individual } X \text{ selected}) = \frac{Eval(X)}{\sum_{Y \in \text{population}} Eval(Y)}$$

- Example (*tournament*): Select a random small subset of the population and select the best-fit individual as a parent
- Implements “survival of the fittest”
- Corresponds loosely to the greedy part of hill-climbing (we try to move uphill)

## GA and Hill Climbing

- Create initial population  $X = \{X_1, \dots, X_P\}$
- Iterate:
  1. Select  $K$  random parents  $(X, X')$ :
 


Hill-climbing component: Try to move uphill as much as possible
  2. For each pair of parents  $(X, X')$ :
    - 1.1 Generate offsprings  $(Y_1, Y_2)$  using crossover operation
 

Random walk component: Move randomly to escape shallow local maxima
    - 1.2 Generate offspring  $Y_i$ :
 

Randomly selected element of the population by  $Y_i$
    - 1.3 With probability  $\mu$ :
 


Apply a random mutation to  $Y_i$
- Return the best individual in the population

## How would you set up these problems to use GA search?

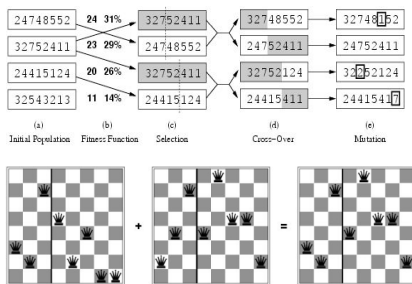


$A \vee \neg B \vee C$   
 $\neg A \vee C \vee D$   
 SAT  $B \vee D \vee \neg E$   
 $\neg C \vee \neg D \vee \neg E$   
 $\neg A \vee \neg C \vee E$   
 .....


N-Queens



## GA for N-queens




## TSP Example

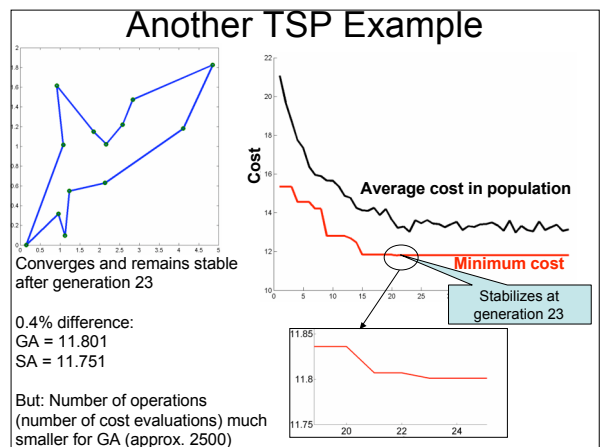





$N = 13$

$P = 100$  elements in population

$\mu = 4\%$  mutation rate  
 $r = 50\%$  reproduction rate  
 $(K = rP)$







## GA Discussion

- Many parameters to tweak:  $\mu$ ,  $P$ ,  $r$
- Many variations on basic scheme. Examples:
  - Multiple-point crossover
  - Dynamic encoding
  - Selection based on rank or relative fitness to least fit individual
  - Multiple fitness functions
  - Combine with a local optimizer (for example, local hill-climbing) → Deviates from "pure" evolutionary view
- In many problems, assuming correct choice of parameters, can be surprisingly effective

## GA Discussion

- Why does it work at all?
- Limited theoretical results (informally!):
  - Suppose that there exists a partial assignment of genes  $s$  such that:
 
$$\text{Average of } Eval(X) \geq \text{Average of } Eval(Y)$$

$X \text{ contains } s$ 
 $Y \in \text{Population}$
  - Then the number of individuals containing  $s$  will increase in the next generation
- Key consequence: The design of the representation (the chromosomes) is critical to the performance the GA. It is probably more important than the choice of parameters of selection strategy, etc.